

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

```
In [2]: df=pd.read_csv(r"C:\Users\pavan\Downloads\data.csv")
df
```

Out[2]:

	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	sqft_above	sqft_basement	yr_built
<b>0</b>	2014-05-02 00:00:00	3.130000e+05	3.0	1.50	1340	7912	1.5	0	0	3	1340	0	1955
<b>1</b>	2014-05-02 00:00:00	2.384000e+06	5.0	2.50	3650	9050	2.0	0	4	5	3370	280	1921
<b>2</b>	2014-05-02 00:00:00	3.420000e+05	3.0	2.00	1930	11947	1.0	0	0	4	1930	0	1966
<b>3</b>	2014-05-02 00:00:00	4.200000e+05	3.0	2.25	2000	8030	1.0	0	0	4	1000	1000	1963
<b>4</b>	2014-05-02 00:00:00	5.500000e+05	4.0	2.50	1940	10500	1.0	0	0	4	1140	800	1976
...	...	...	...	...	...	...	...	...	...	...	...	...	...
<b>4595</b>	2014-07-09 00:00:00	3.081667e+05	3.0	1.75	1510	6360	1.0	0	0	4	1510	0	1954
<b>4596</b>	2014-07-09 00:00:00	5.343333e+05	3.0	2.50	1460	7573	2.0	0	0	3	1460	0	1983
<b>4597</b>	2014-07-09 00:00:00	4.169042e+05	3.0	2.50	3010	7014	2.0	0	0	3	3010	0	2009
<b>4598</b>	2014-07-10 00:00:00	2.034000e+05	4.0	2.00	2090	6630	1.0	0	0	3	1070	1020	1974
<b>4599</b>	2014-07-10 00:00:00	2.206000e+05	3.0	2.50	1490	8102	2.0	0	0	4	1490	0	1990

4600 rows × 18 columns



```
In [4]: df=df[['sqft_living','sqft_above']]  
df.columns=['living','above']
```

```
In [5]: df.head(10)
```

```
Out[5]:
```

	living	above
0	1340	1340
1	3650	3370
2	1930	1930
3	2000	1000
4	1940	1140
5	880	880
6	1350	1350
7	2710	2710
8	2430	1570
9	1520	1520

```
In [6]: df.describe()
```

```
Out[6]:
```

	living	above
count	4600.000000	4600.000000
mean	2139.346957	1827.265435
std	963.206916	862.168977
min	370.000000	370.000000
25%	1460.000000	1190.000000
50%	1980.000000	1590.000000
75%	2620.000000	2300.000000
max	13540.000000	9410.000000

In [7]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4600 entries, 0 to 4599
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0    living   4600 non-null     int64
1    above    4600 non-null     int64
dtypes: int64(2)
memory usage: 72.0 KB
```

In [8]: df.fillna(method='ffill',inplace=True)

C:\Users\pavan\AppData\Local\Temp\ipykernel\_1300\4116506308.py:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

```
df.fillna(method='ffill',inplace=True)
```


In [9]: x=np.array(df['living']).reshape(-1,1)  
y=np.array(df['above']).reshape(-1,1)

In [10]: x\_train,x\_test,y\_train,y\_test=train\_test\_split(x,y,test\_size=0.25)

In [11]: regr=LinearRegression()  
regr.fit(x\_train,y\_train)  
print(regr.score(x\_test,y\_test))

0.7335389803325467

```
In [12]: y_pred=regr.predict(x_test)
plt.scatter(x_test,y_test,color='b')
plt.plot(x_test,y_pred,color='k')
plt.show()
```



```
In [13]: df500=df[:][:500]
```

```
In [16]: df500=df[:][:500]
sns.lmplot(x="living",y="above",data=df500,order=1,ci=None)
df500.fillna(method='ffill',inplace=True)
x=np.array(df500['living']).reshape(-1,1)
y=np.array(df500['above']).reshape(-1,1)
df500.dropna(inplace=True)
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
regr=LinearRegression
regr.fit(x_train,y_train)
print("Regression:",regr.score(x_test,y_test))
y_pred=regr.predict(x_test)
plt.scatter(x_test,y_test,color='b')
plt.plot(x_test,y_pred,color='k')
plt.show()
```


-----  
**TypeError**

Traceback (most recent call last)

Cell In[16], line 9

```
7 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
8 regr=LinearRegression
----> 9 regr.fit(x_train,y_train)
10 print("Regression:",regr.score(x_test,y_test))
11 y_pred=regr.predict(x_test)
```

**TypeError:** LinearRegression.fit() missing 1 required positional argument: 'y'



```
In [17]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
model=LinearRegression()
model.fit(x_train,y_train)
y_pred=model.predict(x_test)
r2=r2_score(y_test,y_pred)
print("R2 score:",r2)
```

R2 score: 0.715532733363012



In [ ]: